

Analysis of Daily COVID-19 Cases, Hospitalizations, and Deaths in NYC

Anja S, Bin Y, Safiya S, Ting-Hsuan C, Yijin W

2023-04-04

Introduction

Let t_i be the number of days since the beginning of a pandemic wave. Let $(y_i, t_i)_{i=1, \dots, n}$ be a sequence of observed daily cases/hospitalizations/deaths at time t_i . Let $Y_i = \sum_{k=1}^i y_k$ be the cumulative number of cases/hospitalizations/deaths by time t_i . Assume that (Y_i, t_i) follows the following non-linear model:

$$Y_i = N(t_i, \boldsymbol{\theta}) + \epsilon_i$$

where $N(t_i, \boldsymbol{\theta})$ is the Richard's growth function with parameters $\boldsymbol{\theta} = (a, k, d, t_0)^T$ and ϵ_i is the random error with mean zero, i.e. $E[\epsilon_i] = 0$.

Task 1

Our goal for Task 1.1 was to develop a Newton-Raphson algorithm to fit Richard curves to each NYC borough's cumulative number of cases for a pandemic wave. We used a simple gradient descent algorithm that aims to minimize the sum of squared errors (SSE), $h(\boldsymbol{\theta})$.

Since $E[Y_i] = N(t_i, \boldsymbol{\theta})$, then $h(\boldsymbol{\theta})$ is defined as:

$$\begin{aligned} h(\boldsymbol{\theta}) &= \sum_{i=1}^n (Y_i - N(t_i, \boldsymbol{\theta}))^2 \\ &= \sum_{i=1}^n \left[Y_i - a \{1 + d \exp \{-k(t_i - t_0)\}\}^{-1/d} \right]^2 \end{aligned}$$

The gradient of $h(\boldsymbol{\theta})$, $\nabla h(\boldsymbol{\theta})$, is calculated as:

$$\begin{aligned} \nabla h(\boldsymbol{\theta}) &= -2 \sum_{i=1}^n [Y_i - N(t_i, \boldsymbol{\theta})] \cdot \nabla N(t_i, \boldsymbol{\theta}) \\ &= -2 \sum_{i=1}^n [Y_i - N(t_i, \boldsymbol{\theta})] \cdot \left(\frac{\partial N(t_i, \boldsymbol{\theta})}{\partial a}, \frac{\partial N(t_i, \boldsymbol{\theta})}{\partial k}, \frac{\partial N(t_i, \boldsymbol{\theta})}{\partial d}, \frac{\partial N(t_i, \boldsymbol{\theta})}{\partial t_0} \right)^T \end{aligned}$$

where

$$\begin{aligned}
\frac{\partial N(t_i, \boldsymbol{\theta})}{\partial a} &= \frac{1}{(1 + de^{-k(t-t_0)})^{1/d}} \\
\frac{\partial N(t_i, \boldsymbol{\theta})}{\partial k} &= \frac{a(t-t_0)e^{-k(t-t_0)}}{(1 + de^{-k(t-t_0)})^{1+1/d}} \\
\frac{\partial N(t_i, \boldsymbol{\theta})}{\partial d} &= -\frac{a(e^{-k(t-t_0)}d - \ln(1 + e^{-k(t-t_0)}d)(1 + e^{-k(t-t_0)}d))}{d^2(1 + de^{-k(t-t_0)})^{1+1/d}} \\
\frac{\partial N(t_i, \boldsymbol{\theta})}{\partial t_0} &= -\frac{ake^{-k(t-t_0)}}{(1 + de^{-k(t-t_0)})^{1+1/d}}
\end{aligned}$$

Our gradient descent algorithm can be summarized as follows:

1. Set $\boldsymbol{\theta}_0$, the starting values for $\boldsymbol{\theta}$ (see later in this section for more details on how to pick these).
2. Update $\boldsymbol{\theta}$ based on $\boldsymbol{\theta}_j = \boldsymbol{\theta}_{j-1} - \lambda \cdot I_{4 \times 4} \nabla h(\boldsymbol{\theta}_{j-1})$, where λ is a user-chosen learning rate (setting λ is to a small value, such as 1^{-10} , works well for this data).
3. If $h(\boldsymbol{\theta}_j) \geq h(\boldsymbol{\theta}_{j-1})$, then decrease the learning rate further and recalculate $h(\boldsymbol{\theta}_j)$, replacing λ with $\frac{\lambda}{10}$. Continue repeating this step until $h(\boldsymbol{\theta}_j) < h(\boldsymbol{\theta}_{j-1})$, i.e. until there is a decrease in SSE from the previous iteration.
4. Continue repeating steps 3 and 4 until convergence is reached, i.e. when the absolute difference between $h(\boldsymbol{\theta}_j)$ and $h(\boldsymbol{\theta}_{j-1})$ is smaller than a very small tolerance level.

Although our algorithm is not the most efficient algorithm, it has two main advantages. One, it is simple to compute and does not rely on the calculation of the Hessian, which is complicated. Two, using the symmetric and positive definite matrix $I_{4 \times 4}$ as a replacement for the Hessian guarantees that we have a descent direction. This means that we will be able to find some $\lambda \in (0, 1)$ that ensures that the updated $\boldsymbol{\theta}_j$ has a smaller SSE than the previous iteration's $\boldsymbol{\theta}_{j-1}$.


In step 1 of the algorithm, the user is required to choose the starting values, $\boldsymbol{\theta}_0$. We found that poor choices of $\boldsymbol{\theta}_0$ resulted in non-convergence or incorrect convergence issues. Choices of $\boldsymbol{\theta}_0$ based on the observed cumulative cases for the borough and the epidemiological interpretations of the Richard growth parameters resulted in fast convergence and good final estimates of Y_i . Our guidelines for how to tailor the starting values for each borough and wave are as follows:

- Since a is the upper bound of the Richard growth function, then let its starting value be the maximum number of cumulative cases in the pandemic wave. For the first wave, this is easily calculated as $\max(Y_i)$. For the second wave, since we only have observed data for the first half of the wave (i.e. up until about the inflection point), we must estimate what the maximum number of cumulative cases will be. We recommend using about $2 \times \max(Y_i)$ under the assumption that half of the maximum cumulative cases occurs at the inflection point. We made this assumption because we saw this trend across the boroughs in the first wave and assumed that the second wave would follow a similar pattern.
- Since t_0 is the time of inflection, then let its starting value be the time t_i where the inflection point occurs. This can be chosen as the point where the curve goes from convex to concave based on a plot of Y_i against t_i . Note that for the second wave, since we only have observed data for the first half of the wave, this point will occur towards the end of the observed data.
- Since k is the growth rate that controls the slope at the inflection point, then let its starting value be calculated as the slope at the inflection point t_0 standardized by the cumulative cases at the inflection point: $\frac{(Y_{i,t_0+m} - Y_{i,t_0-m}) / [(t_0+m) - (t_0-m)]}{Y_{i,t_0}}$, where Y_{i,t_0-m} and Y_{i,t_0+m} are the cumulative cases corresponding to m days before and after the inflection point, respectively, Y_{i,t_0} is the cumulative cases corresponding to the inflection point, and $(t_0+m) - (t_0-m)$ is a small range around the inflection point. This can easily be computed by using two points around the inflection point to fit a linear regression, regressing Y_i on t_i .
- Since d has no epidemiological interpretation, picking its starting value is tricky. We found that values between 0 and 0.5 worked well. However, we advise investigators to try a few different values for d and

pick what works best.

Figure XXX gives an example of how a starting curve for Wave 1 data on NYC would look, compared to the final fitted curve. We can see that the starting curve follows the general trend of the observed cumulative cases, but the final curve obtained from running our algorithm fits the observed data much better. We also note that from about $t_i = 100$ to $t_i = 200$ of the first wave, the observed cumulative cases follow a linear trend, so we would expect that our final Richard growth curve will never be able to fit the observed data very well in this portion of the wave since the Richard growth function assumes a sigmoidal shape.

The algorithm described in this section was also used in Task 1.2 to fit Richard curves to each NYC borough's cumulative number of hospitalizations and deaths for each pandemic wave.



	a	k	d	t0
<i>start</i>	234933	0.062	0.5	35
<i>final</i>	234933	0.04	-0.169	34.976

Figure 1: Example for Wave 1 NYC data comparing starting values and starting curve to the final values and fitted curve obtained from our algorithm.

Task 2

Task 3

Conclusions